

Welcome to Winter 2020 edition of CME 241:

Reinforcement Learning for Stochastic Control Problems in Finance

Instructor: [Ashwin Rao](#)

• **Classes:** Wed & Fri 4:30-5:50pm. [Bldg 380 \(Sloan Mathematics Center - Math Corner\), Room 380w](#)

• **Office Hours:** Fri 2-4pm (or by appointment) in ICME M05 (Huang Engg Bldg)

Overview of the Course

- Theory of Markov Decision Processes (MDPs)
- Dynamic Programming (DP) Algorithms
- Reinforcement Learning (RL) Algorithms
- Plenty of Python implementations of models and algorithms
- We apply these algorithms to 5 Financial/Trading problems:
 - (Dynamic) Asset-Allocation to maximize Utility of Consumption
 - Pricing and Hedging of Derivatives in an Incomplete Market
 - Optimal Exercise/Stopping of Path-dependent American Options
 - Optimal Trade Order Execution (managing Price Impact)
 - Optimal Market-Making (Bid/Ask managing Inventory Risk)
- By treating each of the problems as MDPs (i.e., Stochastic Control)
- We will go over classical/analytical solutions to these problems
- Then we will introduce real-world considerations, and tackle with RL (or DP)
- The course blends Theory/Mathematics, Programming/Algorithms and Real-World Financial Nuances

Learning Material will be a combination of

- [Technical Documents/Lecture Slides I have prepared specifically for this course](#)
- [Python codebase I have developed for this course](#) to help you "learn through coding"
- [Slides and Videos from David Silver's UCL course on RL](#)
- For deeper self-study and reference, augment the above content with [The Sutton-Barto RL Book and Sutton's accompanying teaching material](#)

Lecture-by-Lecture (tentative) schedule with corresponding lecture slides, reading/videos, and assignments

Date	Lecture Slides	Reading/Videos	Suggested Assignments
January	Course Overview		

8		<ul style="list-style-type: none"> • First (Introduction) chapter of Sutton-Barto (pages 1-12) • Optional: Rich Sutton's corresponding slides on Intro to RL • Optional: David Silver's slides on Intro to RL • Optional: David Silver's corresponding video (youtube) on Intro to RL 	<ul style="list-style-type: none"> • Register for the Course on Piazza • Install/Setup on your laptop with LaTeX, Python 3 (and optionally Jupyter notebook) • Create a git repo for this course where you can upload and organize all the code and technical writing you will do as part of assignments and self-learning • Send me a message on Piazza with your git repo URL, so I can periodically review your assignments and other self-learning work
January 10	Markov Processes (MP) and Markov Reward Processes (MRP)	<ul style="list-style-type: none"> • Optional: David Silver's corresponding video (on youtube) on MPs/MRPs/MDPs 	<ul style="list-style-type: none"> • Write out the MP/MRP definitions and MRP Value Function definition (in LaTeX) in your own style/notation (so you really internalize these concepts) • Think about the data structures/class design (in Python 3) to represent MP/MRP and implement them with clear type declarations • Remember - your data structure/code design must resemble the Mathematical/notational formalism as much as possible • Specifically the data structure/code design of MRP should be incremental (and not independent) to that of MP • Separately implement the $r(s,s')$ and the $R(s) = \sum_{s'} p(s,s') * r(s,s')$ definitions of MRP • Write code to convert/cast the $r(s,s')$ definition of MRP to the $R(s)$ definition of MRP

			<p>(put some thought into code design here)</p> <ul style="list-style-type: none"> Write code to generate the stationary distribution for an MDP
January 15	Markov Decision Processes (MDP), Value Function, and Bellman Equations	<ul style="list-style-type: none"> Third (MDP) chapter of Sutton-Barto book (pages 47-67) Optional: Rich Sutton's corresponding slides on MDPs Two ways of arriving at the identical MDP from an MDPRefined ($r(s,s',a)$ definition) 	<ul style="list-style-type: none"> Write the Bellman equation for MDP Value Function and code to calculate MDP Value Function (based on Matrix inversion method you learnt in this lecture) Write out the MDP definition, Policy definition and MDP Value Function definition (in LaTeX) in your own style/notation (so you really internalize these concepts) Think about the data structure/class design (in Python 3) to represent MDP, Policy, Value Function, and implement them with clear type definitions The data structure/code design of MDP should be incremental (and not independent) to that of MRP Separately implement the $r(s,s',a)$ and $R(s,a) = \sum_{s'} p(s,s',a) * r(s,s',a)$ definitions of MDP Write code to convert/cast the $r(s,s',a)$ definition of MDP to the $R(s,a)$ definition of MDP (put some thought into code design here) Write code to create a MRP given a MDP and a Policy Write out all 8 MDP Bellman Equations and also the transformation from Optimal Action-Value function to

			Optimal Policy (in LaTeX)
January 17	Dynamic Programming Algorithms	<ul style="list-style-type: none"> • Fourth (Dynamic Programming) chapter of Sutton-Barto book (pages 73-88) • Understanding Dynamic Programming through Bellman Operators • Optional: Rich Sutton's corresponding slides on Dynamic Programming • Optional: David Silver's video (on youtube) on Dynamic Programming 	Further assignment suggestions will show up here as the course progresses ...
January 22	Understanding Risk-Aversion through Utility Theory (as a pre-req for Finance Applications)	<ul style="list-style-type: none"> • Optional (Related) Reading: A Terse Introduction to Efficient Frontier Mathematics 	
January 24	Application Problem 1 - Optimal Asset Allocation/Consumption (Merton's 1969 Portfolio Problem)	<ul style="list-style-type: none"> • Optional Review: Stochastic Calculus Foundations (used in setting up HJB) • Reference: A paper on A Deep RL Framework for Optimal Asset Allocation • Some (rough) pointers on Discrete versus Continuous MDPs, and solution techniques 	
January 29	Application Problem 2 - Pricing and Hedging of Derivatives in Incomplete Markets and Application Problem 3 - Optimal Exercise of American Options	<ul style="list-style-type: none"> • Optional Review: Foundations of Arbitrage-Free and Complete Markets • Reference: JP Morgan Research paper on Deep Hedging • Reference: Longstaff-Schwartz paper on Pricing American Options (industry-standard approach) • Reference: A paper on RL for Optimal Exercise of American Options 	
January 31	Application Problem 4 - Optimal Trade Order	<ul style="list-style-type: none"> • Reference: Bertsimas-Lo 	

	Execution and Application Problem 5 - Optimal Market-Making	paper on Optimal Trade Order Execution <ul style="list-style-type: none"> • Reference: Almgren-Chriss paper on Risk-Adjusted Optimal Trade Order Execution • Reference: Avellaneda-Stoikov paper on Optimal Market-Making 	
February 5	Industry Guest Lecture		
February 7	Model-free (RL) Prediction With Monte Carlo and Temporal Difference	<ul style="list-style-type: none"> • Optional: David Silver's corresponding video (youtube) on Model-free Prediction • Monte-Carlo and TD (Model-Free) Prediction sections from Sutton-Barto textbook (pages 91-95, 119-128) 	
February 10-11	Take-home Midterm Exam		
February 12	Model-free (RL) Prediction with Eligibility Traces (TD(Lambda))	<ul style="list-style-type: none"> • Optional: David Silver's corresponding video (youtube) on Model-free Prediction • n-Step TD section of Sutton-Barto textbook (pages 141-145) • Optional: TD(Lambda) and Eligibility Traces-based Prediction is covered on pages 287-297, but this treatment is for the more general case of function approximation of Value Function (we've only covered tabular RL algorithms so far). 	
February 14	Model-free Control (RL for Optimal Value Function/Policy)	<ul style="list-style-type: none"> • Optional: David Silver's corresponding video (youtube) on Model-free Control • MC and TD-based Control sections of Sutton-Barto textbook (pages 96-111, 129-134, 146-149) 	

		<ul style="list-style-type: none"> Optional: SARSA(Lambda) is covered on pages 303-307, but this treatment is for the more general case of function approximation of Value Function (we've only covered tabular RL algorithms so far) 	
February 19 and 21	RL with Function Approximation (including Deep RL and Batch Methods)	<ul style="list-style-type: none"> Optional: David Silver's corresponding video (youtube) on RL with Function Approximation Function Approximation sections of Sutton-Barto textbook (pages 197-210, 222-230, 243-248) Optional: Original DQN paper and Nature DQN paper Optional: Lagoudakis-Parr paper on Least Squares Policy Iteration (LSPI) 	
February 26 and 28	Value Function Geometry and Gradient TD		
March 4	Policy Gradient Algorithms	<ul style="list-style-type: none"> Optional: David Silver's corresponding video (youtube) on Policy Gradient Algorithms Policy Gradient chapter of Sutton-Barto textbook (pages 321-332, 335-336) Optional: Original Paper on Policy Gradient 	
March 6	Integrating Learning and Planning	<ul style="list-style-type: none"> Optional: David Silver's corresponding video (youtube) on Integrating Learning and Planning Chapter of Sutton-Barto textbook on Integrating Learning and Planning (pages 159-188) 	
March 11	Exploration versus Exploitation	<ul style="list-style-type: none"> Optional: David Silver's corresponding video (youtube) on Exploration versus Exploitation 	

		<ul style="list-style-type: none"> • Chapter of Sutton-Barto textbook on Multi-Armed Bandits (pages 25-41) 	
March 13	Special Topics: Evolutionary Strategies and Adaptive Multistage Sampling/Monte-Carlo Tree Search Algorithms	<ul style="list-style-type: none"> • Reference: Chang, Fu, Hu, Marcus paper on Adaptive Multistage Sampling • Reference: OpenAI Research paper on Evolutionary Strategies as an Alternative to RL 	<ul style="list-style-type: none"> • Upload all of your assignment work on the github account you had created at the start of the course • Send me a message on Piazza that you have uploaded all of your work • Ensure that all of your assignment work can be accessed by me at github.com (make sure you have pushed and not just committed)
March 17-18	Take-home Final Exam		

Grade will be based on

- 25% Mid-Term Exam (on Theory, Modeling and Algorithms)
- 40% Final Exam (on Theory, Modeling and Algorithms)
- 35% Assignments: Programming, Technical Writing and Theory Problem-Solving (to be done throughout the course)

Purpose and Grading of Assignments

- Assignments are not to be treated as "tests/exams" with a right/wrong answer
- Rather, they should be treated as part of your learning experience
- You will TRULY understand ideas/models/algorithms only when you WRITE down the Mathematics and the Code precisely
- In other words, simply reading the Mathematics or the Code gives you a false sense of understanding things
- Take the initiative to make up your own assignments, especially on topics you feel you don't quite understand
- Individual assignments won't get a grade and there are no due dates for the assignments
- Rather, the entire body of assignments work throughout the course will be graded (upload regularly on your course git repo)
- It will be graded less on correctness and completeness, and more on:
 - Coding and Technical Writing style that is clear and modular
 - Demonstration of curiosity and commitment to learning through the overall body of assignments work
 - Extent of engagement in asking questions and seeking feedback for improvements